

Practical No. 7

Exam Seat No: 22510064 – Parshwa Herwade

Github Link: [Sem-7-Assign/HPC lab at main · parshwa913/Sem-7-Assign · GitHub](#)

1. **Implement Matrix-Vector Multiplication using MPI. Use different number of processes and analyze the performance.**

```
2. #include <stdio.h>
3. #include <stdlib.h>
4. #include <mpi.h>
5.
6. int main(int argc, char* argv[]) {
7.     int rank, size;
8.     int n;
9.     int *matrix = NULL, *vector = NULL, *result = NULL;
10.    int *local_matrix, *local_result;
11.    int rows_per_proc;
12.
13.    MPI_Init(&argc, &argv);
14.    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
15.    MPI_Comm_size(MPI_COMM_WORLD, &size);
16.
17.    if (rank == 0) {
18.        printf("Enter size of square matrix (n): ");
19.        fflush(stdout);
20.        scanf("%d", &n);
21.    }
22.
23.    MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
24.
25.    if (n % size != 0) {
26.        if (rank == 0) {
```

```
27.         printf("Error: n (%d) must be divisible by number of
    processes (%d)\n", n, size);
28.     }
29.     MPI_Finalize();
30.     return 0;
31. }
32.
33. rows_per_proc = n / size;
34.
35. if (rank == 0) {
36.     matrix = (int*)malloc(n * n * sizeof(int));
37.     vector = (int*)malloc(n * sizeof(int));
38.     result = (int*)malloc(n * sizeof(int));
39.
40.     printf("Enter matrix (%d x %d):\n", n, n);
41.     for (int i = 0; i < n; i++)
42.         for (int j = 0; j < n; j++)
43.             scanf("%d", &matrix[i * n + j]);
44.
45.     printf("Enter vector (%d elements):\n", n);
46.     for (int i = 0; i < n; i++)
47.         scanf("%d", &vector[i]);
48. }
49.
50. local_matrix = (int*)malloc(rows_per_proc * n * sizeof(int));
51. local_result = (int*)malloc(rows_per_proc * sizeof(int));
52. if (rank != 0) vector = (int*)malloc(n * sizeof(int));
53.
54. MPI_Scatter(matrix, rows_per_proc * n, MPI_INT,
55.             local_matrix, rows_per_proc * n, MPI_INT,
56.             0, MPI_COMM_WORLD);
57.
58. MPI_Bcast(vector, n, MPI_INT, 0, MPI_COMM_WORLD);
59.
60. for (int i = 0; i < rows_per_proc; i++) {
61.     local_result[i] = 0;
62.     for (int j = 0; j < n; j++) {
63.         local_result[i] += local_matrix[i * n + j] * vector[j];
64.     }
```

```
65.     }
66.
67.     MPI_Gather(local_result, rows_per_proc, MPI_INT,
68.               result, rows_per_proc, MPI_INT,
69.               0, MPI_COMM_WORLD);
70.
71.     if (rank == 0) {
72.         printf("Result vector:\n");
73.         for (int i = 0; i < n; i++)
74.             printf("%d ", result[i]);
75.         printf("\n");
76.     }
77.
78.     if (rank == 0) { free(matrix); free(vector); free(result); }
79.     else free(vector);
80.     free(local_matrix);
81.     free(local_result);
82.
83.     MPI_Finalize();
84.     return 0;
85. }
86.
```

OUTPUT:

```
posh@LAPTOP-ELUQQMKU:~/hpc_assign/assign7$ mpicc matrix_vector.c -o matrix_vector
posh@LAPTOP-ELUQQMKU:~/hpc_assign/assign7$ mpirun -np 4 ./matrix_vector
Enter size of square matrix (n): 3
Error: n (3) must be divisible by number of processes (4)
posh@LAPTOP-ELUQQMKU:~/hpc_assign/assign7$ mpicc matrix_vector.c -o matrix_vector
posh@LAPTOP-ELUQQMKU:~/hpc_assign/assign7$ mpirun -np 4 ./matrix_vector
Enter size of square matrix (n): 2
Error: n (2) must be divisible by number of processes (4)
posh@LAPTOP-ELUQQMKU:~/hpc_assign/assign7$ mpicc matrix_vector.c -o matrix_vector
posh@LAPTOP-ELUQQMKU:~/hpc_assign/assign7$ mpirun -np 2 ./matrix_vector
Enter size of square matrix (n): 4
Enter matrix (4 x 4):
1 2 3 4
2 3 4 5
3 4 5 6
4 5 6 7
Enter vector (4 elements):
5 6 7 8
Result vector:
70 96 122 148
posh@LAPTOP-ELUQQMKU:~/hpc_assign/assign7$ mpirun -np 2 ./matrix_vector
Enter size of square matrix (n): 2
Enter matrix (2 x 2):
1 2
3 4
Enter vector (2 elements):
4 5
Result vector:
14 32
posh@LAPTOP-ELUQQMKU:~/hpc_assign/assign7$
```

Algorithm

1. Initialize the MPI environment.
2. Process 0 (root) takes the size n , the matrix A , and the vector x as input.
3. The rows of matrix A are divided among the processes (block row distribution).
 - o Each process gets n/p rows (if n divisible by p).
4. The vector x is broadcast to all processes.
5. Each process computes its partial product:

$$y_i = \sum_{j=0}^{n-1} A_{ij} \cdot x_j$$

for its assigned rows.

6. The partial results are gathered at the root process using MPI_Gather.
7. Root process prints the result vector.
8. Finalize MPI.

Observations (Sample Outputs)

- Execution time decreases as the number of processes increases (for large matrices).
- For small n, communication overhead may dominate, giving no real speedup.

Conclusion

- Matrix–vector multiplication parallelizes well because rows can be distributed independently.
- Speedup is noticeable for larger matrices.
- For small matrices, MPI overhead reduces efficiency.

2. Implement Matrix-Matrix Multiplication using MPI. Use different number of processes and analyze the performance.

```
3. #include <stdio.h>
4. #include <stdlib.h>
5. #include <mpi.h>
```

```
6.
7. int main(int argc, char* argv[]) {
8.     int rank, size;
9.     int n;
10.    int *A = NULL, *B = NULL, *C = NULL;
11.    int *local_A, *local_C;
12.    int rows_per_proc;
13.
14.    MPI_Init(&argc, &argv);
15.    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
16.    MPI_Comm_size(MPI_COMM_WORLD, &size);
17.
18.    if (rank == 0) {
19.        printf("Enter size of square matrices (n): ");
20.        fflush(stdout);
21.        scanf("%d", &n);
22.    }
23.
24.    MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
25.
26.    if (n % size != 0) {
27.        if (rank == 0) {
28.            printf("Error: n (%d) must be divisible by number of
processes (%d)\n", n, size);
29.        }
30.        MPI_Finalize();
31.        return 0;
32.    }
33.
34.    rows_per_proc = n / size;
35.
36.    if (rank == 0) {
37.        A = (int*)malloc(n * n * sizeof(int));
38.        B = (int*)malloc(n * n * sizeof(int));
39.        C = (int*)malloc(n * n * sizeof(int));
40.
41.        printf("Enter matrix A (%d x %d):\n", n, n);
42.        for (int i = 0; i < n; i++)
43.            for (int j = 0; j < n; j++)
```

```
44.         scanf("%d", &A[i * n + j]);
45.
46.         printf("Enter matrix B (%d x %d):\n", n, n);
47.         for (int i = 0; i < n; i++)
48.             for (int j = 0; j < n; j++)
49.                 scanf("%d", &B[i * n + j]);
50.     }
51.
52.     local_A = (int*)malloc(rows_per_proc * n * sizeof(int));
53.     local_C = (int*)malloc(rows_per_proc * n * sizeof(int));
54.     if (rank != 0) B = (int*)malloc(n * n * sizeof(int));
55.
56.     MPI_Scatter(A, rows_per_proc * n, MPI_INT,
57.                local_A, rows_per_proc * n, MPI_INT,
58.                0, MPI_COMM_WORLD);
59.
60.     MPI_Bcast(B, n * n, MPI_INT, 0, MPI_COMM_WORLD);
61.
62.     for (int i = 0; i < rows_per_proc; i++) {
63.         for (int j = 0; j < n; j++) {
64.             local_C[i * n + j] = 0;
65.             for (int k = 0; k < n; k++) {
66.                 local_C[i * n + j] += local_A[i * n + k] * B[k *
n + j];
67.             }
68.         }
69.     }
70.
71.     MPI_Gather(local_C, rows_per_proc * n, MPI_INT,
72.                C, rows_per_proc * n, MPI_INT,
73.                0, MPI_COMM_WORLD);
74.
75.     if (rank == 0) {
76.         printf("Result matrix C:\n");
77.         for (int i = 0; i < n; i++) {
78.             for (int j = 0; j < n; j++)
79.                 printf("%d ", C[i * n + j]);
80.             printf("\n");
81.         }
```

```
82.     }  
83.  
84.     if (rank == 0) { free(A); free(B); free(C); }  
85.     else free(B);  
86.     free(local_A);  
87.     free(local_C);  
88.  
89.     MPI_Finalize();  
90.     return 0;  
91. }  
92.
```

OUTPUT:

```
● posh@LAPTOP-ELUQQMKU:~/hpc_assign/assign7$ mpicc matrix_matrix.c -o matrix_matrix  
● posh@LAPTOP-ELUQQMKU:~/hpc_assign/assign7$ mpirun -np 2 ./matrix_matrix  
Enter size of square matrices (n): 2  
Enter matrix A (2 x 2):  
1 2  
3 4  
Enter matrix B (2 x 2):  
5 6  
7 8  
Result matrix C:  
19 22  
43 50  
● posh@LAPTOP-ELUQQMKU:~/hpc_assign/assign7$ mpirun -np 2 ./matrix_matrix  
Enter size of square matrices (n): 4  
Enter matrix A (4 x 4):  
1 2 3 4  
2 3 4 5  
5 6 7 8  
6 7 8 9  
Enter matrix B (4 x 4):  
1 0 0 0  
0 1 0 0  
0 0 1 0  
0 0 0 1  
Result matrix C:  
1 2 3 4  
2 3 4 5  
5 6 7 8  
6 7 8 9  
○ posh@LAPTOP-ELUQQMKU:~/hpc_assign/assign7$
```


Algorithm

1. Initialize the MPI environment.
2. Process 0 (root) takes size n, and matrices A and B as input.
3. The rows of matrix A are scattered among all processes.
4. Matrix B is broadcast to all processes.
5. Each process computes partial product:

$$C_{ij} = \sum_{k=0}^{n-1} A_{ik} \cdot B_{kj}$$

1. for its assigned rows.
2. Partial results are gathered back at the root process.
3. Root process prints the result matrix.
4. Finalize MPI.

Observations (Sample Outputs)

- Speedup increases with larger matrices but communication overhead affects small cases.

Conclusion

- Matrix-matrix multiplication is highly parallelizable, as computations for rows can be distributed.
- MPI provides good scalability for large n.
- Communication and gathering steps are bottlenecks when n is small.